

SVM Example

The goal is to:

- Learn how to manipulate a SVM in R with the package kernlab
- Observe the effect of changing the C parameter and the kernel
- Test a SVM classifier for cancer diagnosis from gene expression data

1 Linear SVM

Here we generate a toy dataset in 2D, and learn how to train and test a SVM.

1.1 Generate toy data

First generate a set of positive and negative examples from 2 Gaussians.


```
set.seed(1)

n <- 150 #number of data points
p <- 2 # dimension
sigma <- 1 # variance of the distribution
meanpos <- 0 # centre of the distribution of positive examples
meanneg <- 3 # centre of the distribution of negative examples
npos <- round(n / 2) # number of positive examples
nneg <- n - npos # number of negative examples

# Generate the positive and negative examples
xpos <- matrix(rnorm(npos * p, mean = meanpos, sd = sigma), npos, p)
xneg <- matrix(rnorm(nneg * p, mean = meanneg, sd = sigma), nneg, p)
x <- rbind(xpos, xneg)

# Generate the labels
y <- matrix(c(rep(1, npos), rep(-1, nneg)))

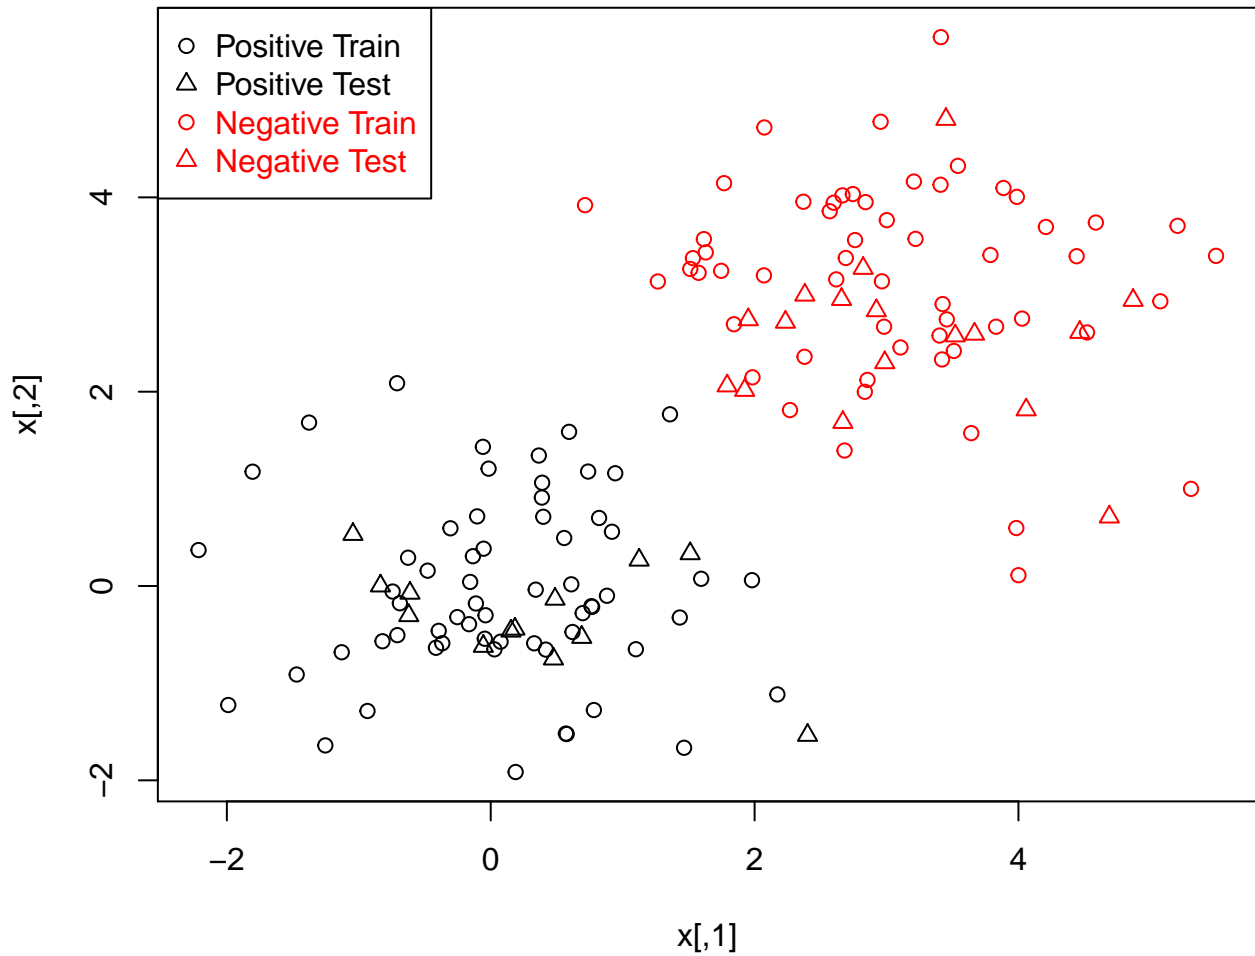
# Visualize the data
plot(x, col = ifelse(y > 0, 1, 2))
legend("topleft", c("Positive", "Negative"), col = seq(2), pch = 1, text.col = seq(2))
```



Now we split the data into a training set (80%) and a test set (20%)

```
# Prepare a training and a test set
ntrain <- round(n * 0.8) # number of training examples
tindex <- sample(n, ntrain) # indices of training samples
xtrain <- x[tindex, ]
xtest <- x[-tindex, ]
ytrain <- y[tindex]
ytest <- y[-tindex]
istrain <- rep(0, n)
istrain[tindex] <- 1

# Visualize
plot(x, col = ifelse(y > 0, 1, 2), pch = ifelse(istrain == 1, 1, 2))
legend("topleft", c("Positive Train", "Positive Test", "Negative Train", "Negative Test"), col = c(1, 1, 2, 2), pch = c(1, 2, 1, 2))
```



1.2 Train a SVM

Now we train a linear SVM with parameter $C=100$ on the training set.

```
# load the kernlab package
# install.packages("kernlab")
library(kernlab)

# train the SVM
svp <- ksvm(xtrain, ytrain, type = "C-svc", kernel = "vanilladot", C=100, scaled=c())
```

```
## Setting default kernel parameters
```

```
#Look and understand what svp contains
# General summary
svp
```

```
## Support Vector Machine object of class "ksvm"
##
## SV type: C-svc (classification)
## parameter : cost C = 100
##
## Linear (vanilla) kernel function.
##
```

```

## Number of Support Vectors : 3
##
## Objective Function Value : -3.7546
## Training error : 0
# Attributes that you can access
attributes(svp)

## $param
## $param$C
## [1] 100
##
##
## $scaling
## `001NULL\001`
##
## $coef
## $coef[[1]]
## [1] -3.5699332  3.7543303 -0.1843971
##
##
## $alphaindex
## $alphaindex[[1]]
## [1]  1 65 90
##
##
## $b
## [1] -6.859059
##
## $obj
## [1] -3.754606
##
## $SVindex
## [1]  1 65 90
##
## $nSV
## [1] 3
##
## $prior
## $prior[[1]]
## $prior[[1]]$prior1
## [1] 62
##
## $prior[[1]]$prior0
## [1] 58
##
##
##
## $prob.model
## $prob.model[[1]]
## NULL
##
##
## $alpha
## $alpha[[1]]

```

```

## [1] 3.5699332 3.7543303 0.1843971
##
##
## $type
## [1] "C-svc"
##
## $kernel
## function (x, y = NULL)
## {
##     if (!is(x, "vector"))
##         stop("x must be a vector")
##     if (!is(y, "vector") && !is.null(y))
##         stop("y must be a vector")
##     if (is(x, "vector") && is.null(y)) {
##         crossprod(x)
##     }
##     if (is(x, "vector") && is(y, "vector")) {
##         if (!length(x) == length(y))
##             stop("number of dimension must be the same on both data points")
##         crossprod(x, y)
##     }
## }
## <environment: 0x438ec38>
## attr("kpar")
## list()
## attr("class")
## [1] "vanillakernel"
## attr("class")attr("package")
## [1] "kernlab"
##
## $kpar
## list()
##
## $xmatrix
## $xmatrix[[1]]
##      X1      X2
## 1  1.984153 2.148143
## 65 1.358680 1.767287
## 90 1.266782 3.134448
##
##
## $ymatrix
## [1] -1 -1  1 -1 -1 -1  1  1 -1  1 -1  1  1 -1 -1  1  1  1 -1  1 -1  1  1
## [24]  1  1 -1  1 -1  1 -1 -1  1  1 -1  1  1  1 -1 -1  1 -1  1 -1 -1 -1  1
## [47]  1 -1 -1 -1  1 -1  1  1 -1  1  1 -1 -1 -1  1  1  1 -1  1  1 -1  1 -1
## [70] -1  1  1  1 -1 -1 -1  1 -1 -1  1  1 -1  1  1  1 -1  1 -1  1 -1  1  1
## [93]  1  1 -1 -1  1  1 -1  1  1 -1 -1 -1  1 -1 -1  1  1 -1  1  1 -1 -1 -1
## [116] -1 -1 -1  1  1
##
## $fitted
## [1] -1 -1  1 -1 -1 -1  1  1 -1  1 -1  1  1 -1 -1  1  1  1 -1  1 -1  1  1
## [24]  1  1 -1  1 -1  1 -1 -1  1  1 -1  1  1  1 -1 -1  1 -1  1 -1 -1 -1  1
## [47]  1 -1 -1 -1  1 -1  1  1 -1  1  1 -1 -1 -1  1  1  1 -1  1  1 -1  1 -1
## [70] -1  1  1  1 -1 -1 -1  1 -1 -1  1  1 -1  1  1  1 -1  1 -1  1 -1  1  1

```

```

## [93]  1  1 -1 -1  1  1 -1  1  1 -1 -1 -1  1 -1 -1  1  1 -1  1  1 -1 -1 -1
## [116] -1 -1 -1  1  1
##
## $lev
## [1] -1  1
##
## $nclass
## [1] 2
##
## $error
## [1] 0
##
## $cross
## [1] -1
##
## $n.action
## function (object, ...)
## UseMethod("na.omit")
## <bytecode: 0x3101a28>
## <environment: namespace:stats>
##
## $terms
## `001NULL\001`
##
## $kcall
## .local(x = x, y = ..1, scaled = ..5, type = "C-svc", kernel = "vanilladot",
##      C = 100)
##
## $class
## [1] "ksvm"
## attr("package")
## [1] "kernlab"

# For example, the support vectors
alpha(svp)

## [[1]]
## [1] 3.5699332 3.7543303 0.1843971

alphaindex(svp)

## [[1]]
## [1] 1 65 90


b(svp)

## [1] -6.859059

# Use the built-in function to pretty-plot the classifier
plot(svp, data = xtrain)

```


SVM classification plot



Next, we create a function `plotlinearsvm=function(svp,xtrain)` to plot the points and the decision boundaries of a linear SVM.

```
plotlinearsvm <- function(svp, xtrain, plot_legend=T){
  plot(xtrain, pch = ifelse(ytrain > 0, 1, 2), xlim = c(-2, 6), ylim = c(-3, 6))
  if (plot_legend)
    legend("topleft", c("Positive", "Negative"), pch = seq(2))
  w = colSums(unlist(alpha(svp)) * ytrain[unlist(alphaindex(svp))] * xtrain[unlist(alphaindex(svp)),])
  b = - b(svp)
  abline(a= -b / w[2], b = -w[1]/w[2])
  abline(a= (-b+1)/ w[2], b = -w[1]/w[2], lty = 2)
  abline(a= (-b-1)/ w[2], b = -w[1]/w[2], lty = 2)
}

plotlinearsvm(svp, xtrain)
```



The figure represents a linear SVM with decision boundary $f(x) = 0$. Dotted lines correspond to the level $f(x) = 1$ and $f(x) = -1$.

1.3 Predict with a SVM

Now we can use the trained SVM to predict the label of points in the test set, and we analyze the results using variant metrics.

```
# Predict labels on test
ypred <- predict(svp, xtest)
table(ytest, ypred)

##      ypred
## ytest -1  1
##      -1 17  0
##       1  0 13

# Compute accuracy
sum(ypred == ytest) / length(ytest)

## [1] 1

# Compute at the prediction scores
ypredscore <- predict(svp, xtest, type = "decision")
```



```
# Check that the predicted labels are the signs of the scores
table(ypredscore > 0, ypred)
```

```
##      ypred
##      -1  1
## FALSE 17  0
## TRUE   0 13
```

```
# Package to compute ROC curve, precision-recall etc...
# install.packages("ROCR")
library(ROCR)
```

```
## Loading required package: gplots
```

```
##
```

```
## Attaching package: 'gplots'
```

```
## The following object is masked from 'package:stats':
```

```
##
```


```
##      lowess
```

```
pred <- prediction(ypredscore, ytest)
```


```
# Plot ROC curve
```

```
perf <- performance(pred, measure = "tpr", x.measure = "fpr")
```


```
plot(perf)
```



```
# Plot precision/recall curve  
perf <- performance(pred, measure = "prec", x.measure = "rec")  
plot(perf)
```



```
# Plot accuracy as function of threshold  
perf <- performance(pred, measure = "acc")  
plot(perf)
```



2 Cross-validation

Instead of fixing a training set and a test set, we can improve the quality of these estimates by running k-fold cross-validation. We split the training set in k groups of approximately the same size, then iteratively train a SVM using k - 1 groups and make prediction on the group which was left aside. When k is equal to the number of training points, we talk of leave-one-out (LOO) cross-validation. To generate a random split of n points in k folds, we can for example create the following function:

```
cv.folds <- function(y, folds = 3){
  ## randomly split the n samples into folds
  split(sample(length(y)), rep(1:folds, length = length(y)))
}
```

We write a function `cv.ksvm = function(x, y, folds = 3,...)` which returns a vector `ypred` of predicted decision score for all points by k-fold cross-validation.

```
cv.ksvm <- function(x, y, folds = 3,...){
  index = cv.folds(y, folds = folds)
  predScore = rep(NA, length(y))
  for (i in 1:folds){
    toTrain = unname(unlist(index[-i]))
    testSet = index[[i]]
    svp = ksvm(x[toTrain, ], y[toTrain], type = "C-svc",
```

```

        kernel = "vanilladot", C=100, scaled=c())
    predScore[testSet] = predict(svp, x[unlist(index[[i]]), ], type = "decision")
  }
  predScore
}

```

We compute the various performance of the SVM by 5-fold cross-validation.

```
ypredscore = cv.ksvm(x, y, folds=5)
```

```


## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters

```

```

pred = prediction(ypredscore, y)
perf = performance(pred, measure = "tpr", x.measure = "fpr")
plot(perf)


```



```

perf = performance(pred, measure = "acc")
plot(perf)

```



Alternatively, the `ksvm` function can automatically compute the k-fold cross-validation accuracy.

```
svp <- ksvm(x, y, type = "C-svc", kernel = "vanilladot", C = 100, scaled=c(), cross = 5)
```

```
## Setting default kernel parameters
```

```
print(cross(svp))
```

```
## [1] 0.006666667
```

```
print(error(svp))
```

```
## [1] 0
```

2.1 Effect of C


The `C` parameter balances the trade-off between having a large margin and separating the positive and unlabeled on the training set. It is important to choose it well to have good generalization.

Plot the decision functions of SVM trained on the toy examples for different values of `C` in the range $2^{\text{seq}(-10, 14)}$.

```
cost = 2^(seq(-10, 14, by=3))
par(mfrow = c(3,3))
for (c in cost){
  svp = ksvm(xtrain, ytrain, type = "C-svc", kernel = "vanilladot", C=c, scaled=c())
```

```
plotlinearsvm(svp, xtrain, plot_legend = F)
}
```


```
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
```




```
par(mfrow=c(1,1))
```

Plot the 5-fold cross-validation error as a function of C .

```
cost = 2^(seq(-10, 15))
crossError = rep(NA, length(cost))
```

```
plot(log2(cost), error, type='b')
```


Do the same on data with more overlap between the two classes, e.g., re-generate toy data with `meanneg` being 1.

```
n <- 150 #number of data points
p <- 2 # dimension
sigma <- 1 # variance of the distribution
meanpos <- 0 # centre of the distribution of positive examples
meanneg <- 1 # centre of the distribution of negative examples
npos <- round(n / 2) # number of positive examples
nneg <- n - npos # number of negative examples

# Generate the positive and negative examples
xpos <- matrix(rnorm(npos * p, mean = meanpos, sd = sigma), npos, p)
xneg <- matrix(rnorm(nneg * p, mean = meanneg, sd = sigma), nneg, p)
x <- rbind(xpos, xneg)

# Generate the labels
y <- matrix(c(rep(1, npos), rep(-1, nneg)))

# Visualize the data
plot(x, col = ifelse(y > 0, 1, 2))
legend("topleft", c("Positive", "Negative"), col = seq(2), pch = 1, text.col = seq(2))
```




```
# generate training/testing
ntrain <- round(n * 0.8) # number of training examples
tindex <- sample(n, ntrain) # indices of training samples
xtrain <- x[tindex, ]
xtest <- x[-tindex, ]
ytrain <- y[tindex]
ytest <- y[-tindex]
istrain <- rep(0, n)
istrain[tindex] <- 1

# cost cross validation
cost = 2^(seq(-10, 15))
error = sapply(cost, function(c){
  cross(ksvm(x, y, type = "C-svc", kernel = "vanilladot", C = c, scaled=c(),
    cross = 5))
})


## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
```

```
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
## Setting default kernel parameters
```

```
plot(cost, error, type='b')
```



```
plot(log2(cost), error, type='b')
```




3 Nonlinear SVM

Sometimes linear SVM are not enough. For example, generate a toy dataset where positive and negative examples are mixture of two Gaussians which are not linearly separable.

```
RandomMatrix <- function( dist, n, p, ... ) {
  rs <- dist( n*p, ... )
  matrix( rs, n, p )
}

GenerateDatasetNonlinear <- function( n, p ) {
  bottom.left <- RandomMatrix( rnorm, n, p, mean=0, sd=1 )
  upper.right <- RandomMatrix( rnorm, n, p, mean=4, sd=1 )
  tmp1 <- RandomMatrix( rnorm, n, p, mean=0, sd=1 )
  tmp2 <- RandomMatrix( rnorm, n, p, mean=4, sd=1 )
  upper.left <- cbind( tmp1[,1], tmp2[,2] )
  bottom.right <- cbind( tmp2[,1], tmp1[,2] )
  y <- c( rep( 1, 2 * n ), rep( -1, 2 * n ) )
  idx.train <- sample( 4 * n, floor( 3.5 * n ) )
  is.train <- rep( 0, 4 * n )
  is.train[idx.train] <- 1
  data.frame( x=rbind( bottom.left, upper.right, upper.left, bottom.right ), y=y, train=is.train )
}

data = GenerateDatasetNonlinear(150, 2)
plot(data[,1:2], col = data[,3] + 2)
```




```
x = as.matrix(data[,1:2])
y = matrix(data[,3])
```

To solve this problem, we should instead use a nonlinear SVM. This is obtained by simply changing the kernel parameter. For example, to use a Gaussian RBF kernel with $\sigma = 1$ and $C = 1$:

```
# Train a nonlinear SVM
svp <- ksvm(x, y, type = "C-svc", kernel="rbf", kpar = list(sigma = 1), C = 1)

# Visualize it
plot(svp, data = x)
```

SVM classification plot



You should obtain something that look like Figure 3. Much better than the linear SVM, no? The nonlinear SVM has now two parameters: σ and C . Both play a role in the generalization capacity of the SVM.

Visualize and compute the 5-fold cross-validation error for different values of C and σ . Observe their influence.

```
library(ggplot2)


##
## Attaching package: 'ggplot2'
## The following object is masked from 'package:kernlab':
##
##   alpha
cost = 2^(seq(-10, 15, by=2))
sigma = 0:5
error = sapply(cost, function(c){
  sapply(sigma, function(s){
    cross(ksvm(x, y, type = "C-svc", kernel="rbf", kpar = list(sigma = s), C = c,
      scaled=c(), cross = 5))
  })
})
toPlotError = data.frame(sigma = rep(sigma, length(cost)),
```

```

cost = rep(cost, each = length(sigma)),
error = as.vector(error))

ggplot(data = toPlotError, aes(x=cost, y=error)) + geom_point() + geom_line() +
  facet_grid(.~sigma)

```



A useful heuristic to choose σ is implemented in kernlab. It is based on the quantiles of the distances between the training point.


```

# Train a nonlinear SVM with automatic selection of sigma by heuristic
svp <- ksvm(x, y, type = "C-svc", kernel = "rbf", C = 1)

# Visualize it
plot(svp, data = x)


```


SVM classification plot



Train a non-linear SVM with various of C with automatic determination of σ . In fact, many other non-linear kernels are implemented. Check the documentation of kernlab to see them: `?kernels`

```
library(ggplot2)
cost = 2^(seq(-10, 15, by=2))
error = sapply(1:length(cost), function(i){
  svp = ksvm(x, y, type = "C-svc", kernel = "rbf", C = cost[i], cross = 5)
  cross(svp)
})
plot(cost, error, type="o")
```




Test the polynomial, hyperbolic tangent, Laplacian, Bessel and ANOVA kernels on the toy examples.

```
myKernels = c("polydot", "tanhdot", "laplacedot", "besseldot", "anovadot")
for (kernel in myKernels){
  plot(ksvm(x, y, type = "C-svc", kernel = kernel, C = 1), data=x)
}
```


```
## Setting default kernel parameters
```

SVM classification plot




```
## Setting default kernel parameters
```

SVM classification plot




SVM classification plot




```
## Setting default kernel parameters
```

SVM classification plot



Setting default kernel parameters

SVM classification plot



4 Application: cancer diagnosis from gene expression data

As a real-world application, let us test the ability of SVM to predict the class of a tumour from gene expression data. We use a publicly available dataset of gene expression data for 128 different individuals with acute lymphoblastic leukemia (ALL).

```
# To install this package run
#
# source("https://bioconductor.org/biocLite.R")
# biocLite("ALL")
#
#
# Load the ALL dataset
library(ALL)

## Loading required package: Biobase
## Loading required package: BiocGenerics
## Loading required package: parallel
##
## Attaching package: 'BiocGenerics'
## The following objects are masked from 'package:parallel':
##
##   clusterApply, clusterApplyLB, clusterCall, clusterEvalQ,
##   clusterExport, clusterMap, parApply, parCapply, parLapply,
##   parLapplyLB, parRapply, parSapply, parSapplyLB
## The following objects are masked from 'package:stats':
##
##   IQR, mad, xtabs
## The following objects are masked from 'package:base':
##
##   anyDuplicated, append, as.data.frame, cbind, colnames,
##   do.call, duplicated, eval, evalq, Filter, Find, get, grep,
##   grepl, intersect, is.unsorted, lapply, lengths, Map, mapply,
##   match, mget, order, paste, pmax, pmax.int, pmin, pmin.int,
##   Position, rank, rbind, Reduce, rownames, sapply, setdiff,
##   sort, table, tapply, union, unique, unsplit, which, which.max,
##   which.min
## Welcome to Bioconductor
##
##   Vignettes contain introductory material; view with
##   'browseVignettes()'. To cite Bioconductor, see
##   'citation("Biobase")', and for packages 'citation("pkgname)".

data(ALL)

# Inspect them
?ALL
show(ALL)

## ExpressionSet (storageMode: lockedEnvironment)
## assayData: 12625 features, 128 samples
##   element names: exprs
```



```
## protocolData: none
## phenoData
##   sampleNames: 01005 01010 ... LAL4 (128 total)
##   varLabels: cod diagnosis ... date last seen (21 total)
##   varMetadata: labelDescription
## featureData: none
## experimentData: use 'experimentData(object)'
##   pubMedIds: 14684422 16243790
## Annotation: hgu95av2

print(summary(pData(ALL)))

##      cod      diagnosis      sex      age
## Length:128      Length:128      F :42      Min. : 5.00
## Class :character      Class :character      M :83      1st Qu.:19.00
## Mode :character      Mode :character      NA's: 3      Median :29.00
##                                     Mean :32.37
##                                     3rd Qu.:45.50
##                                     Max. :58.00
##                                     NA's :5
##      BT      remission      CR      date.cr
## B2 :36      CR :99      Length:128      Length:128
## B3 :23      REF :15      Class :character      Class :character
## B1 :19      NA's:14      Mode :character      Mode :character
## T2 :15
## B4 :12
## T3 :10
## (Other):13
## t(4;11)      t(9;22)      cyto.normal      citog
## Mode :logical      Mode :logical      Mode :logical      Length:128
## FALSE:86      FALSE:67      FALSE:69      Class :character
## TRUE :7      TRUE :26      TRUE :24      Mode :character
## NA's :35      NA's :35      NA's :35
##
##
##      mol.biol      fusion protein      mdr      kinet      ccr
## ALL1/AF4:10      p190 :17      NEG :101      dyploid:94      Mode :logical
## BCR/ABL :37      p190/p210: 8      POS : 24      hyperd.:27      FALSE:74
## E2A/PBX1: 5      p210 : 8      NA's: 3      NA's : 7      TRUE :26
## NEG :74      NA's :95      NA's :28
## NUP-98 : 1
## p15/p16 : 1
##
## relapse      transplant      f.u      date last seen
## Mode :logical      Mode :logical      Length:128      Length:128
## FALSE:35      FALSE:91      Class :character      Class :character
## TRUE :65      TRUE :9      Mode :character      Mode :character
## NA's :28      NA's :28
##
##
##
```

Here we focus on predicting the type of the disease (B-cell or T-cell). We get the expression data and disease type as follows

```
x <- t(exprs(ALL))
y <- substr(ALL$BT,1,1)
```

Test the ability of a SVM to predict the class of the disease from gene expression. Check the influence of the parameters.

You actually need to play a lot with the parameters: kernels, costs, ...

```
x <- t(exprs(ALL))
y <- substr(ALL$BT,1,1)

# train and test sets
n = length(y)
ntrain <- round(n * 0.8) # number of training examples
tindex <- sample(n, ntrain) # indices of training samples
xtrain <- x[tindex, ]
xtest <- x[-tindex, ]
ytrain <- y[tindex]
ytest <- y[-tindex]

# train svm on train set
svp = ksvm(xtrain, ytrain, type = "C-svc", kernel = "rbf", C = 10)

# predict on test set
(pred = predict(svp, xtest))
```

```
## [1] B B B B B B B B B B B B B B B B T T B T B T T T
## Levels: B T
```

```
(acc = sum(as.vector(pred) == ytest) / length(ytest))
```

```
## [1] 0.9230769
```

Finally, we may want to predict the type and stage of the diseases. We are then confronted with a multi-class classification problem, since the variable to predict can take more than two values:

```
y <- ALL$BT
print(y)
```

```
## [1] B2 B2 B4 B1 B2 B1 B1 B1 B2 B2 B3 B3 B3 B2 B3 B B2 B3 B2 B3 B2 B2 B2
## [24] B1 B1 B2 B1 B2 B1 B2 B B B2 B2 B2 B1 B2 B2 B2 B2 B2 B4 B4 B2 B2 B2
## [47] B4 B2 B1 B2 B2 B3 B4 B3 B3 B3 B4 B3 B3 B1 B1 B1 B1 B3 B3 B3 B3 B3 B3
## [70] B3 B3 B1 B3 B1 B4 B2 B2 B1 B3 B4 B4 B2 B2 B3 B4 B4 B4 B1 B2 B2 B2 B1
## [93] B2 B B T T3 T2 T2 T3 T2 T T4 T2 T3 T3 T T2 T3 T2 T2 T2 T1 T4 T
## [116] T2 T3 T2 T2 T2 T2 T3 T3 T3 T2 T3 T2 T
## Levels: B B1 B2 B3 B4 T T1 T2 T3 T4
```

Fortunately, kernlab implements automatically multi-class SVM by an all-versus-all strategy to combine several binary SVM.

Test the ability of a SVM to predict the class and the stage of the disease from gene expression.

```
# now y is multiclass
x <- t(exprs(ALL))
y <- ALL$BT

# train and test sets (same as previous question)
n = length(y)
```

```

ntrain <- round(n * 0.8) # number of training examples
tindex <- sample(n, ntrain) # indices of training samples
xtrain <- x[tindex, ]
xtest <- x[-tindex, ]
ytrain <- y[tindex]
ytest <- y[-tindex]

# train svm on train set
# type of svm is now able to handle multiclass
svp = ksvm(xtrain, ytrain, type = "kbb-svc", kernel = "rbf", C = 10)

# predict on test set
(pred = predict(svp, xtest))

## [1] T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1 T1
## [24] T1 T1 T1
## Levels: B B1 B2 B3 B4 T T1 T2 T3 T4

(acc = sum(as.vector(pred) == ytest) / length(ytest))

## [1] 0.03846154

```